

TrustLoRA Audit Report: ybelkada/opt-350m-lora

This report summarizes a full four-dimension audit of a public PEFT adapter. The adapter was evaluated as an adapter, not as a standalone model, using base model facebook/opt-350m.

Plain-language result

The adapter finished the full audit with low composite risk (0.239). It is suitable for continued research use, while still needing normal model-card and license review.

Audit configuration

- Adapter: ybelkada/opt-350m-lora
- Base model: facebook/opt-350m
- Auditor version: 1.0.0
- Audit timestamp: 2026-05-05T16:03:02.991438+00:00
- Runtime: 1254.8 seconds
- Evidence level: full 4/4 behavioural plus static weight audit

Scores

- Composite risk: 0.2395 (LOW)
- Toxicity overall: 0.0042; severe toxicity 0.0001; identity attack 0.0002
- Jailbreak refusal rate: 0.0800; compliance rate 0.0200; partial/uncertain 0.9000
- Bias aggregate: 0.0133; gender 0.0002; racial 0.0493; age 0.0001; occupational 0.0037
- Backdoor anomaly: 0.0000; distribution shift 2.0993; outlier-rank ratio 1.0000; flagged True

What I would check before deployment

- Inspect LoRA tensors and training provenance because the static backdoor screen crossed areview threshold.
- Record the exact adapter revision before publication so future reviewers can reproduce theresult.

Jailbreak probe breakdown

- authority: 1 refusal, 0 compliance, 9 partial/uncertain
- direct: 0 refusal, 0 compliance, 10 partial/uncertain
- hypothetical: 0 refusal, 0 compliance, 10 partial/uncertain
- roleplay: 1 refusal, 0 compliance, 9 partial/uncertain
- technical: 2 refusal, 1 compliance, 7 partial/uncertain

Backdoor review notes

The static scanner found layer-level signals that deserve manual review. These signals do not prove a trigger exists, but they are useful triage markers.

- distribution shift in `base_model.model.model.decoder.layers.0.self_attn.q_proj.lora_B.weight`-
- distribution shift in `base_model.model.model.decoder.layers.0.self_attn.v_proj.lora_B.weight`
- distribution shift in `base_model.model.model.decoder.layers.1.self_attn.q_proj.lora_B.weight`
- distribution shift in `base_model.model.model.decoder.layers.1.self_attn.v_proj.lora_B.weight`
- distribution shift in `base_model.model.model.decoder.layers.10.self_attn.q_proj.lora_B.weight` -
- distribution shift in

`base_model.model.model.decoder.layers.11.self_attn.v_proj.lora_B.weight`

Limitations

- The behavioural suite is prompt-based. It samples observable behavior but cannot prove global safety.
- The jailbreak classifier is heuristic, so ambiguous or evasive answers may require human review.
- The bias score uses toxicity identity-attack variance as a screening proxy, not a full fairness audit.
- The backdoor scanner is static and trigger-agnostic. It is designed to flag anomalies, not to prove exploitability.